

What is Pandas?

- Pandas is a Python library used for working with data sets.
- It has functions for analyzing, cleaning, exploring, and manipulating data.
- The name "Pandas" has a reference to both "Panel Data", and "Python Data Analysis" and was created by Wes McKinney in 2008.

Why Use Pandas?

- Pandas allows us to analyze big data and make conclusions based on statistical theories.
- Pandas can clean messy data sets, and make them readable and relevant.
- Relevant data is very important in data science.

Installation of Pandas

```
In [1]: # pip install pandas
```

Import Pandas as pd

Pandas is usually imported under the pd alias.

- alias: In Python alias are an alternate name for referring to the same thing.

```
In [3]: import pandas as pd
# From pandas import *
# From pandas import pd
```

What is a Series?

- A Pandas Series is like a column in a table.
- It is a one-dimensional array holding data of any type.

```
In [10]: import pandas as pd

data = pd.Series([0.25, 0.5, 0.75, 1.0]) #A Pandas Series is a one-dimensional array of indexed
data
```

```
Out[10]: 0    0.25
1    0.50
2    0.75
3    1.00
dtype: float64
```

Labels

- If nothing else is specified, the values are labeled with their index number. First value has index 0, second value has index 1 etc.

- This label can be used to access a specified value.

```
In [11]: print( data[0])
```

```
0.25
```

Create Labels

- With the (index) argument, you can name your own labels.

```
In [12]: data.index
```

```
Out[12]: RangeIndex(start=0, stop=4, step=1)
```

```
In [13]: data = pd.Series([0.25, 0.5, 0.75, 1.0],index=['a', 'b', 'c', 'd'])
data
```

```
Out[13]: a    0.25
b    0.50
c    0.75
d    1.00
dtype: float64
```

```
In [15]: print( data['b'])
```

```
0.5
```

```
In [16]: data.values
```

```
Out[16]: array([0.25, 0.5 , 0.75, 1.  ])
```

```
In [17]: data[:]
```

```
Out[17]: a    0.25
b    0.50
c    0.75
d    1.00
dtype: float64
```

```
In [18]: data[:2]
```

```
Out[18]: a    0.25
b    0.50
dtype: float64
```

```
In [19]: data[2:]
```

```
Out[19]: c    0.75
d    1.00
dtype: float64
```

Series-as-dictionary

```
In [20]: #Series-as-dictionary
population_dict = {'California': 38332521,
                  'Texas': 26448193,
                  'New York': 19651127,
```

```
'Florida': 19552860,  
'Illinois': 12882135}
```

```
population = pd.Series(population_dict)  
population
```

```
Out[20]: California    38332521  
Texas      26448193  
New York   19651127  
Florida    19552860  
Illinois    12882135  
dtype: int64
```

```
In [21]: population['California']
```

```
Out[21]: 38332521
```

```
In [24]: population[:, 'New York']
```

```
Out[24]: California    38332521  
Texas      26448193  
New York   19651127  
dtype: int64
```

DataFrames

DataFrame is an analog of a two-dimensional array with both flexible row indices and flexible column names

A DataFrame is a collection of Series objects, and a singlecolumn DataFrame can be constructed from a single Series

```
In [31]: import pandas as pd  
  
data = {"calories": [420, 380, 390],  
        "duration": [50, 40, 45]}  
  
#Load data into a DataFrame object  
df = pd.DataFrame(data)  
  
print(df)
```

```
   calories  duration  
0        420        50  
1        380        40  
2        390        45
```

```
In [32]: df.index #index labels
```

```
Out[32]: RangeIndex(start=0, stop=3, step=1)
```

```
In [34]: df.columns #column labels
```

```
Out[34]: Index(['calories', 'duration'], dtype='object')
```

```
In [35]: df['calories']
```

```
Out[35]: 0    420  
1    380  
2    390  
Name: calories, dtype: int64
```

Locate Row

As you can see from the result above, the DataFrame is like a table with rows and columns.

Pandas use the loc attribute to return one or more specified row(s)

```
In [38]: # #refer to the row index:
print(df.loc[0])
```

```
calories    420
duration     50
Name: 0, dtype: int64
```

```
In [39]: # Example : Return row 0 and 1:
#use a list of indexes:
print(df.loc[[0, 1]])
```

```
   calories  duration
0        420        50
1        380        40
```

```
In [40]: # Named Indexes
import pandas as pd

data = {"calories": [420, 380, 390],
        "duration": [50, 40, 45]}

df = pd.DataFrame(data, index = ["day1", "day2", "day3"])

print(df)
```

```
   calories  duration
day1      420        50
day2      380        40
day3      390        45
```

```
In [41]: print(df.loc["day2"])
```

```
calories    380
duration     40
Name: day2, dtype: int64
```

Extracting and transforming data (Load Files Into a DataFrame)

```
In [43]: # Read file .CSV
df1 = pd.read_csv('F:\\Job\\FCDS\\Data Science Methodology\\Data sets\\Online Retail.csv')
df1
```

```
Out[43]:
```

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	12/1/2010 8:26	2.55	17850.0	United Kingdom
1	536365	71053	WHITE METAL LANTERN	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	12/1/2010 8:26	2.75	17850.0	United Kingdom

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
...
541904	581587	22613	PACK OF 20 SPACEBOY NAPKINS	12	12/9/2011 12:50	0.85	12680.0	France
541905	581587	22899	CHILDREN'S APRON DOLLY GIRL	6	12/9/2011 12:50	2.10	12680.0	France
541906	581587	23254	CHILDRENS CUTLERY DOLLY GIRL	4	12/9/2011 12:50	4.15	12680.0	France
541907	581587	23255	CHILDRENS CUTLERY CIRCUS PARADE	4	12/9/2011 12:50	4.15	12680.0	France
541908	581587	22138	BAKING SET 9 PIECE RETROSPOT	3	12/9/2011 12:50	4.95	12680.0	France

541909 rows × 8 columns

In [45]:

```
# Read in filename and set the index: election (index_col='column name')
df = pd.read_csv('F:\\Job\\FCDS\\Data Science Methodology\\Data sets\\Online Retail.csv',
                 index_col='InvoiceNo')
df
```

Out[45]:

	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
InvoiceNo							
536365	85123A	WHITE HANGING HEART T- LIGHT HOLDER	6	12/1/2010 8:26	2.55	17850.0	United Kingdom
536365	71053	WHITE METAL LANTERN	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	12/1/2010 8:26	2.75	17850.0	United Kingdom
536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
...
581587	22613	PACK OF 20 SPACEBOY NAPKINS	12	12/9/2011 12:50	0.85	12680.0	France
581587	22899	CHILDREN'S APRON DOLLY GIRL	6	12/9/2011 12:50	2.10	12680.0	France
581587	23254	CHILDRENS CUTLERY DOLLY GIRL	4	12/9/2011 12:50	4.15	12680.0	France
581587	23255	CHILDRENS CUTLERY CIRCUS PARADE	4	12/9/2011 12:50	4.15	12680.0	France
581587	22138	BAKING SET 9 PIECE RETROSPOT	3	12/9/2011 12:50	4.95	12680.0	France

541909 rows × 7 columns

```
In [46]: # Indexing using square brackets
# data name['column name']['row name']
df['Description']['536365']
```

```
Out[46]: 536365    WHITE HANGING HEART T-LIGHT HOLDER
536365              WHITE METAL LANTERN
536365    CREAM CUPID HEARTS COAT HANGER
536365    KNITTED UNION FLAG HOT WATER BOTTLE
536365    RED WOOLLY HOTTIE WHITE HEART.
536365    SET 7 BABUSHKA NESTING BOXES
536365    GLASS STAR FROSTED T-LIGHT HOLDER
Name: Description, dtype: object
```

```
In [47]: # Using column attribute and row label
df.Description['581587']
```

```
Out[47]: 581587    CIRCUS PARADE LUNCH BOX
581587    PLASTERS IN TIN CIRCUS PARADE
581587    PLASTERS IN TIN STRONGMAN
581587    ALARM CLOCK BAKELIKE PINK
581587    ALARM CLOCK BAKELIKE RED
581587    ALARM CLOCK BAKELIKE GREEN
581587    ALARM CLOCK BAKELIKE IVORY
581587    CHILDRENS APRON SPACEBOY DESIGN
581587    SPACEBOY LUNCH BOX
581587    CHILDRENS CUTLERY SPACEBOY
581587    PACK OF 20 SPACEBOY NAPKINS
581587    CHILDREN'S APRON DOLLY GIRL
581587    CHILDRENS CUTLERY DOLLY GIRL
581587    CHILDRENS CUTLERY CIRCUS PARADE
581587    BAKING SET 9 PIECE RETROSPOT
Name: Description, dtype: object
```

Pandas - Analyzing DataFrames

The head() method returns the headers and a specified number of rows, starting from the top.

```
In [61]: # if the number of rows is not specified, the head() method will return the top 5 rows.
print(df.head())
```

	StockCode	Description	Quantity	\
InvoiceNo				
536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	
536365	71053	WHITE METAL LANTERN	6	
536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	
536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	
536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	

	InvoiceDate	UnitPrice	CustomerID	Country
InvoiceNo				
536365	12/1/2010 8:26	2.55	17850.0	United Kingdom
536365	12/1/2010 8:26	3.39	17850.0	United Kingdom
536365	12/1/2010 8:26	2.75	17850.0	United Kingdom
536365	12/1/2010 8:26	3.39	17850.0	United Kingdom
536365	12/1/2010 8:26	3.39	17850.0	United Kingdom

```
In [62]: print(df.head(10))
```

	StockCode	Description	Quantity	\
InvoiceNo				
536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	
536365	71053	WHITE METAL LANTERN	6	
536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	
536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	
536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	
536365	22752	SET 7 BABUSHKA NESTING BOXES	2	
536365	21730	GLASS STAR FROSTED T-LIGHT HOLDER	6	
536366	22633	HAND WARMER UNION JACK	6	

536366	22632	HAND WARMER RED POLKA DOT	6
536367	84879	ASSORTED COLOUR BIRD ORNAMENT	32

	InvoiceDate	UnitPrice	CustomerID	Country
InvoiceNo				
536365	12/1/2010 8:26	2.55	17850.0	United Kingdom
536365	12/1/2010 8:26	3.39	17850.0	United Kingdom
536365	12/1/2010 8:26	2.75	17850.0	United Kingdom
536365	12/1/2010 8:26	3.39	17850.0	United Kingdom
536365	12/1/2010 8:26	3.39	17850.0	United Kingdom
536365	12/1/2010 8:26	7.65	17850.0	United Kingdom
536365	12/1/2010 8:26	4.25	17850.0	United Kingdom
536366	12/1/2010 8:28	1.85	17850.0	United Kingdom
536366	12/1/2010 8:28	1.85	17850.0	United Kingdom
536367	12/1/2010 8:34	1.69	13047.0	United Kingdom

The tail() method returns the headers and a specified number of rows, starting from the bottom.

```
In [63]: # Print the last 5 rows of the DataFrame
print(df.tail())
```

	StockCode	Description	Quantity \
InvoiceNo			
581587	22613	PACK OF 20 SPACEBOY NAPKINS	12
581587	22899	CHILDREN'S APRON DOLLY GIRL	6
581587	23254	CHILDRENS CUTLERY DOLLY GIRL	4
581587	23255	CHILDRENS CUTLERY CIRCUS PARADE	4
581587	22138	BAKING SET 9 PIECE RETROSPOT	3

	InvoiceDate	UnitPrice	CustomerID	Country
InvoiceNo				
581587	12/9/2011 12:50	0.85	12680.0	France
581587	12/9/2011 12:50	2.10	12680.0	France
581587	12/9/2011 12:50	4.15	12680.0	France
581587	12/9/2011 12:50	4.15	12680.0	France
581587	12/9/2011 12:50	4.95	12680.0	France

The DataFrames object has a method called info(), that gives you more information about the data set.

```
In [64]: print(df.info())

<class 'pandas.core.frame.DataFrame'>
Index: 541909 entries, 536365 to 581587
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   StockCode       541909 non-null object
1   Description      540455 non-null object
2   Quantity        541909 non-null int64
3   InvoiceDate      541909 non-null object
4   UnitPrice       541909 non-null float64
5   CustomerID      406829 non-null float64
6   Country         541909 non-null object
dtypes: float64(2), int64(1), object(4)
memory usage: 53.1+ MB
None
```

```
In [66]: print(df.describe())
```

	Quantity	UnitPrice	CustomerID
count	541909.000000	541909.000000	406829.000000
mean	9.552250	4.611114	15287.690570
std	218.081158	96.759853	1713.600303
min	-80995.000000	-11062.060000	12346.000000
25%	1.000000	1.250000	13953.000000
50%	3.000000	2.080000	15152.000000
75%	10.000000	4.130000	16791.000000
max	80995.000000	38970.000000	18287.000000

merge datasets

In [116...]

```
df1 = pd.DataFrame({'employee': ['Bob', 'Jake', 'Lisa', 'Sue'],
                    'group': ['Accounting', 'Engineering', 'Engineering', 'HR']})
df2 = pd.DataFrame({'employee': ['Lisa', 'Bob', 'Jake', 'Sue'],
                    'hire_date': [2004, 2008, 2012, 2014]})

print(df1)
print(df2)
```

```
   employee    group
0       Bob  Accounting
1       Jake  Engineering
2       Lisa  Engineering
3        Sue         HR

   employee  hire_date
0       Lisa      2004
1        Bob      2008
2       Jake      2012
3        Sue      2014
```

In [117...]

```
df3 = pd.merge(df1, df2)
df3
```

Out[117...]

	employee	group	hire_date
0	Bob	Accounting	2008
1	Jake	Engineering	2012
2	Lisa	Engineering	2004
3	Sue	HR	2014

Accessors

A more efficient and more programmatically reusable method of accessing data in a DataFrame is by using accessors

- loc - accesses using labels
- iloc - accesses using index positions

Both accessors use left bracket, row specifier, comma, column specifier, right bracket as syntax

In [67]:

```
# Using the .loc accessor
df.loc['581587', 'Country']
```

Out[67]:

```
InvoiceNo
581587    France
581587    France
581587    France
581587    France
581587    France
581587    France
581587    France
581587    France
581587    France
581587    France
581587    France
581587    France
581587    France
581587    France
581587    France
Name: Country, dtype: object
```

In [68]:

```
# Using the .iloc accessor
```



```
df.iloc[4, 6]
```

```
Out[68]: 'United Kingdom'
```

Selecting only some columns

- When using bracket-indexing without the `.loc` or `.iloc` accessors, the result returned can be an individual value, Pandas Series, or Pandas DataFrame.
- To ensure the return value is a DataFrame, use a nested list within square brackets

```
In [69]: df_new = df[['CustomerID', 'Country']]
df_new
```

```
Out[69]:
```

	CustomerID	Country
InvoiceNo		
536365	17850.0	United Kingdom
536365	17850.0	United Kingdom
536365	17850.0	United Kingdom
536365	17850.0	United Kingdom
536365	17850.0	United Kingdom
...
581587	12680.0	France
581587	12680.0	France
581587	12680.0	France
581587	12680.0	France
581587	12680.0	France

541909 rows × 2 columns

```
In [70]: df['Country']
```

```
Out[70]: InvoiceNo
536365    United Kingdom
536365    United Kingdom
536365    United Kingdom
536365    United Kingdom
536365    United Kingdom
...
581587    France
581587    France
581587    France
581587    France
581587    France
Name: Country, Length: 541909, dtype: object
```

```
In [71]: type(df.Country)
```

```
Out[71]: pandas.core.series.Series
```

```
In [72]: df['Country'][1:4] # Part of the Country column
```

Out[72]: InvoiceNo
536365 United Kingdom
536365 United Kingdom
536365 United Kingdom
Name: Country, dtype: object

```
In [73]: df['Country'][4] # The value associated with 536365
```

Out[73]: 'United Kingdom'

```
In [74]: df.loc[:, 'CustomerID': 'Country'] # ALL rows, some columns
```

Out[74]:

	CustomerID	Country
InvoiceNo		
536365	17850.0	United Kingdom
536365	17850.0	United Kingdom
536365	17850.0	United Kingdom
536365	17850.0	United Kingdom
536365	17850.0	United Kingdom
...
581587	12680.0	France
581587	12680.0	France
581587	12680.0	France
581587	12680.0	France
581587	12680.0	France

541909 rows × 2 columns

```
In [75]: df.loc['536365':'581587',:] # Some rows, all columns
```

Out[75]:

	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
InvoiceNo							
536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	12/1/2010 8:26	2.55	17850.0	United Kingdom
536365	71053	WHITE METAL LANTERN	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	12/1/2010 8:26	2.75	17850.0	United Kingdom
536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	12/1/2010 8:26	3.39	17850.0	United Kingdom
...
581587	22613	PACK OF 20 SPACEBOY NAPKINS	12	12/9/2011 12:50	0.85	12680.0	France
581587	22899	CHILDREN'S APRON DOLLY GIRL	6	12/9/2011 12:50	2.10	12680.0	France

	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
InvoiceNo							
581587	23254	CHILDRENS CUTLERY DOLLY GIRL	4	12/9/2011 12:50	4.15	12680.0	France
581587	23255	CHILDRENS CUTLERY CIRCUS PARADE	4	12/9/2011 12:50	4.15	12680.0	France
581587	22138	BAKING SET 9 PIECE RETROSPOT	3	12/9/2011 12:50	4.95	12680.0	France

541909 rows × 7 columns

Filtering DataFrames

In [76]: *#Creating a Boolean Series*
df.UnitPrice > 60

Out[76]: InvoiceNo
536365 False
536365 False
536365 False
536365 False
536365 False
...
581587 False
581587 False
581587 False
581587 False
581587 False
Name: UnitPrice, Length: 541909, dtype: bool

In [77]: df[df.UnitPrice > 60]

	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
InvoiceNo							
536392	22827	RUSTIC SEVENTEEN DRAWER SIDEBOARD	1	12/1/2010 10:29	165.00	13705.0	United Kingdom
536544	DOT	DOTCOM POSTAGE	1	12/1/2010 14:32	569.77	NaN	United Kingdom
536592	DOT	DOTCOM POSTAGE	1	12/1/2010 17:06	607.49	NaN	United Kingdom
536676	21769	VINTAGE POST OFFICE CABINET	1	12/2/2010 12:18	79.95	16752.0	United Kingdom
536835	22655	VINTAGE RED KITCHEN CABINET	1	12/2/2010 18:06	295.00	13145.0	United Kingdom
...
581238	DOT	DOTCOM POSTAGE	1	12/8/2011 10:53	1683.75	NaN	United Kingdom
581439	DOT	DOTCOM POSTAGE	1	12/8/2011 16:30	938.59	NaN	United Kingdom
581492	DOT	DOTCOM POSTAGE	1	12/9/2011 10:03	933.17	NaN	United Kingdom
581498	DOT	DOTCOM POSTAGE	1	12/9/2011 10:26	1714.17	NaN	United Kingdom

1188 rows × 7 columns

```
x = df.UnitPrice > 60
df[x]
```

1188 rows \times 7 columns

```
df[(df.UnitPrice >= 50) & (df.CustomerID < 200)] # Both conditions
```

```
df[(df.UnitPrice >= 50) | (df.CustomerID < 200)] # Either condition
```

```
Out[80]:
```

StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
InvoiceNo						

	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
InvoiceNo							
536392	22827	RUSTIC SEVENTEEN DRAWER SIDEBOARD	1	12/1/2010 10:29	165.00	13705.0	United Kingdom
536540	C2	CARRIAGE	1	12/1/2010 14:05	50.00	14911.0	EIRE
536544	22769	CHALKBOARD KITCHEN ORGANISER	1	12/1/2010 14:32	51.02	NaN	United Kingdom
536544	DOT	DOTCOM POSTAGE	1	12/1/2010 14:32	569.77	NaN	United Kingdom
536592	22503	CABIN BAG VINTAGE PAISLEY	1	12/1/2010 17:06	59.53	NaN	United Kingdom
...
581238	DOT	DOTCOM POSTAGE	1	12/8/2011 10:53	1683.75	NaN	United Kingdom
581439	DOT	DOTCOM POSTAGE	1	12/8/2011 16:30	938.59	NaN	United Kingdom
581492	DOT	DOTCOM POSTAGE	1	12/9/2011 10:03	933.17	NaN	United Kingdom
581498	DOT	DOTCOM POSTAGE	1	12/9/2011 10:26	1714.17	NaN	United Kingdom
C581499	M	Manual	-1	12/9/2011 10:28	224.69	15498.0	United Kingdom

1417 rows × 7 columns

Creating a Series

A Pandas Series is a one-dimensional array of indexed data (one column)

```
In [81]: data = { 'Name':pd.Series(['Huda','Mohamed','Hossam','Mina']),
                 'Age': pd.Series([25,30,50,36])}
data
```

```
Out[81]: {'Name': 0      Huda
          1    Mohamed
          2     Hossam
          3       Mina
          dtype: object,
          'Age': 0      25
                 1      30
                 2      50
                 3      36
          dtype: int64}
```

```
In [82]: prices = [10.70, 10.86, 10.74, 10.71, 10.79]
shares = pd.Series(prices)
shares
```

```
Out[82]: 0      10.70
          1      10.86
          2      10.74
          3      10.71
          4      10.79
          dtype: float64
```

```
In [83]: days = ['Mon', 'Tue', 'Wed', 'Thur', 'Fri']
        shares = pd.Series(prices, index=days)
        shares
```

```
Out[83]: Mon      10.70
        Tue      10.86
        Wed      10.74
        Thur     10.71
        Fri      10.79
        dtype: float64
```

Examining an index

```
In [84]: print(shares.index)
        print(shares.index[2])
        print(shares.index[:2])
        print(shares.index[-2:])
        print(shares.index.name)
```

```
Index(['Mon', 'Tue', 'Wed', 'Thur', 'Fri'], dtype='object')
Wed
Index(['Mon', 'Tue'], dtype='object')
Index(['Thur', 'Fri'], dtype='object')
None
```

```
In [85]: # Modifying index name
        shares.index.name = 'weekday'
        shares
```

```
Out[85]: weekday
        Mon      10.70
        Tue      10.86
        Wed      10.74
        Thur     10.71
        Fri      10.79
        dtype: float64
```

```
In [86]: # Modifying all index entries
        shares.index = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday']
        shares
```

```
Out[86]: Monday      10.70
        Tuesday      10.86
        Wednesday    10.74
        Thursday     10.71
        Friday       10.79
        dtype: float64
```

```
In [ ]:
```

```
In [87]: stocks = pd.DataFrame([[ '2016-10-03', 31.50, 14070500, 'CSCO'],
                                [ '2016-10-03', 112.52, 21701800, 'AAPL'],
                                [ '2016-10-03', 57.42, 19189500, 'MSFT'],
                                [ '2016-10-04', 113.00, 29736800, 'AAPL'],
                                [ '2016-10-04', 57.24, 20085900, 'MSFT'],
                                [ '2016-10-04', 31.35, 18460400, 'CSCO'],
                                [ '2016-10-05', 57.64, 16726400, 'MSFT'],
                                [ '2016-10-05', 31.59, 11808600, 'CSCO'],
                                [ '2016-10-05', 113.05, 21453100, 'AAPL']],
                                columns=[ 'Date', 'Close', 'Volume', 'Symbol'])

        stocks
```

```
Out[87]:
```

Date	Close	Volume	Symbol
------	-------	--------	--------

	Date	Close	Volume	Symbol
0	2016-10-03	31.50	14070500	CSCO
1	2016-10-03	112.52	21701800	AAPL
2	2016-10-03	57.42	19189500	MSFT
3	2016-10-04	113.00	29736800	AAPL
4	2016-10-04	57.24	20085900	MSFT
5	2016-10-04	31.35	18460400	CSCO
6	2016-10-05	57.64	16726400	MSFT
7	2016-10-05	31.59	11808600	CSCO
8	2016-10-05	113.05	21453100	AAPL

```
In [88]: # Setting index
stocks = stocks.set_index(['Symbol', 'Date'])
stocks
```

```
Out[88]:
```

	Symbol	Date	Close	Volume
	CSCO	2016-10-03	31.50	14070500
	AAPL	2016-10-03	112.52	21701800
	MSFT	2016-10-03	57.42	19189500
	AAPL	2016-10-04	113.00	29736800
	MSFT	2016-10-04	57.24	20085900
	CSCO	2016-10-04	31.35	18460400
	MSFT	2016-10-05	57.64	16726400
	CSCO	2016-10-05	31.59	11808600
	AAPL	2016-10-05	113.05	21453100

```
In [89]: # MultiIndex on DataFrame
stocks.index
```

```
Out[89]: MultiIndex([(('CSCO', '2016-10-03'),
                        ('AAPL', '2016-10-03'),
                        ('MSFT', '2016-10-03'),
                        ('AAPL', '2016-10-04'),
                        ('MSFT', '2016-10-04'),
                        ('CSCO', '2016-10-04'),
                        ('MSFT', '2016-10-05'),
                        ('CSCO', '2016-10-05'),
                        ('AAPL', '2016-10-05'))],
                    names=['Symbol', 'Date'])
```

```
In [90]: print(stocks.index.name)
```

None

```
In [91]: print(stocks.index.names)
```

['Symbol', 'Date']

```
In [92]:
```

```
# Sorting index
stocks = stocks.sort_index()
stocks
```

```
Out[92]:
```

		Close	Volume
Symbol	Date		
AAPL	2016-10-03	112.52	21701800
	2016-10-04	113.00	29736800
	2016-10-05	113.05	21453100
CSCO	2016-10-03	31.50	14070500
	2016-10-04	31.35	18460400
	2016-10-05	31.59	11808600
MSFT	2016-10-03	57.42	19189500
	2016-10-04	57.24	20085900
	2016-10-05	57.64	16726400

```
In [93]:
```

```
# Indexing (individual row)
stocks.loc[('CSCO', '2016-10-04')]
```

```
Out[93]:
```

Close	31.35
Volume	18460400.00

Name: (CSCO, 2016-10-04), dtype: float64

```
In [94]:
```

```
stocks.loc[('CSCO', '2016-10-04'), 'Volume']
```

```
Out[94]:
```

18460400.0

```
In [95]:
```

```
stocks.loc['AAPL']
```

```
Out[95]:
```

	Close	Volume
Date		
2016-10-03	112.52	21701800
2016-10-04	113.00	29736800
2016-10-05	113.05	21453100

```
In [96]:
```

```
# Slicing (outermost index)
stocks.loc['CSCO':'MSFT']
```

```
Out[96]:
```

		Close	Volume
Symbol	Date		
CSCO	2016-10-03	31.50	14070500
	2016-10-04	31.35	18460400
	2016-10-05	31.59	11808600
MSFT	2016-10-03	57.42	19189500
	2016-10-04	57.24	20085900
	2016-10-05	57.64	16726400

In []:

Stacking & unstacking DataFrames

In [97]:

```
trials = pd.DataFrame([[1, 'A', 'F', 5],
                        [2, 'A', 'M', 3],
                        [3, 'B', 'F', 8],
                        [4, 'B', 'M', 9]],
                        columns=['id', 'treatment', 'gender', 'response'])

trials
```

Out[97]:

	id	treatment	gender	response
0	1	A	F	5
1	2	A	M	3
2	3	B	F	8
3	4	B	M	9

In [98]:

```
trials = trials.set_index(['treatment', 'gender'])
trials
```

Out[98]:

		id	response
treatment		gender	
A	F	1	5
	M	2	3
B	F	3	8
	M	4	9

In [99]:

```
# Unstacking a multi-index
trials.unstack(level='gender')
```

Out[99]:

		id		response	
gender	F	M	F	M	
	treatment				
A	1	2	5	3	
B	3	4	8	9	

In [100]:

```
# Stacking DataFrames
trials_by_gender = trials.unstack(level='gender')
trials_by_gender
```

Out[100]:

		id		response	
gender	F	M	F	M	
	treatment				
A	1	2	5	3	

		id		response	
gender		F	M	F	M
treatment					
B		3	4	8	9

```
In [101... trials_by_gender.stack(level='gender')
```

Out[101...

		id		response	
treatment	gender				
A	F	1		5	
	M	2		3	
B	F	3		8	
	M	4		9	

```
In [102... stacked = trials_by_gender.stack(level='gender')
stacked
```

Out[102...

		id		response	
treatment	gender				
A	F	1		5	
	M	2		3	
B	F	3		8	
	M	4		9	

```
In [103... # Swapping levels
swapped = stacked.swaplevel(0, 1)
swapped
```

Out[103...

		id		response	
gender	treatment				
F	A	1		5	
M	A	2		3	
F	B	3		8	
M	B	4		9	

```
In [104... # Sorting rows
sorted_trials = swapped.sort_index()
sorted_trials
```

Out[104...

		id		response	
gender	treatment				
F	A	1		5	
	B	3		8	

		id	response
gender	treatment		
M	A	2	3
	B	4	9

Categorical and groupby

```
In [105... sales = pd.DataFrame(
    {
        'weekday': ['Sun', 'Sun', 'Mon', 'Mon'],
        'city': ['Austin', 'Dallas', 'Austin', 'Dallas'],
        'bread': [139, 237, 326, 456],
        'butter': [20, 45, 70, 98]
    }
)
sales
```

```
Out[105... weekday  city  bread  butter
0      Sun  Austin   139     20
1      Sun  Dallas   237     45
2     Mon  Austin   326     70
3     Mon  Dallas   456     98
```

```
In [106... # filter and count
sales.loc[sales['weekday'] == 'Sun'].count()
```

```
Out[106... weekday    2
city         2
bread        2
butter       2
dtype: int64
```

```
In [107... # Groupby and count
sales.groupby('weekday').count()
```

```
Out[107... weekday
city  bread  butter
Mon    2     2     2
Sun    2     2     2
```

Aggregation/Reduction

Some reducing functions

- mean()
- std()
- sum()
- first(), last()
- min(), max()

```
In [108... # Groupby and sum
sales.groupby('weekday')['bread'].sum()
```

```
Out[108... weekday
Mon      782
Sun      376
Name: bread, dtype: int64
```

```
In [109... # Groupby and sum: multiple columns
sales
```

```
Out[109...   weekday  city  bread  butter
0      Sun  Austin   139     20
1      Sun  Dallas   237     45
2      Mon  Austin   326     70
3      Mon  Dallas   456     98
```

```
In [110... sales.groupby('weekday')[['bread', 'butter']].sum()
```

```
Out[110...   bread  butter
weekday
Mon      782    168
Sun      376     65
```

Groupby and aggregation

The `.agg()` method can be used with a tuple or list of aggregations as input. When applying multiple aggregations on multiple columns, the aggregated DataFrame has a multi-level column index.

```
In [111... sales
```

```
Out[111...   weekday  city  bread  butter
0      Sun  Austin   139     20
1      Sun  Dallas   237     45
2      Mon  Austin   326     70
3      Mon  Dallas   456     98
```

```
In [112... sales.groupby('city')[['bread', 'butter']].max()
```

```
Out[112...   bread  butter
city
Austin   326     70
Dallas   456     98
```

```
In [113... sales.groupby('city')[['bread', 'butter']].agg(['max', 'sum'])
```

Out[113...

	bread		butter	
	max	sum	max	sum
city				
Austin	326	465	70	90
Dallas	456	693	98	143

Aggregation functions

string names

- 'sum'
- 'mean'
- 'count'

In [114...

```
# Custom aggregation
def data_range(series):
    return series.max() - series.min()
```

In [115...

```
sales.groupby('weekday')[['bread', 'butter']].agg(data_range)
```

Out[115...

	bread	butter
weekday		
Mon	130	28
Sun	98	25

In []: